

# A private intelligence layer built around Level's operating context.

Level Intelligence OS is a Level-specific construction intelligence and pursuit system. It finds project signals earlier, qualifies them against Level's own operating criteria, routes work to the right person, coordinates RFQs and go/no-go decisions, preserves relationship memory, and gives leadership a weekly view of what matters — all on top of an owned operating context that compounds every time the team makes a call.

---

## PREPARED FOR

Level Construction • Founders + BD +  
Precon + Capital

## DOCUMENT

Detailed Proposition • v2026.05

## STATUS

Client-ready draft • for working review

## PREPARED BY

Omar Sayed • Grounded

## COMPANIONS

Two Paths walkthrough • Agent Ops live  
demo

## CONFIDENTIALITY

Shared in confidence with Level  
Construction

## CONTENTS

|    |                                                                           |      |
|----|---------------------------------------------------------------------------|------|
| 01 | Executive summary                                                         | § 1  |
| 02 | The context advantage — and why nothing on the market does this for Level | § 2  |
| 03 | Where Mercator fits, and where Level Intelligence OS goes further         | § 3  |
| 04 | The agentic stack — six specialists that change the workflow game         | § 4  |
| 05 | ROI — eight compounding value levers                                      | § 5  |
| 06 | Procore bridge strategy — wrap first, replace only when earned            | § 6  |
| 07 | Human control and trust                                                   | § 7  |
| 08 | The 90-day pilot                                                          | § 8  |
| 09 | Commercial principles and ownership                                       | § 9  |
| 10 | Discovery questions                                                       | § 10 |
| 11 | Closing position                                                          | § 11 |
| 12 | Companion artifacts                                                       | § 12 |

# 1 • Executive summary

Level already has the hard part: trusted relationships, commercial construction judgment, development instincts, hospitality experience, capital discipline, and a practical sense for which opportunities are worth time. The opportunity is to turn that judgment into a private operating system that Level owns.

**The thesis is simple.** Generic SaaS makes Level adapt to vendor screens. Level Intelligence OS does the inverse: the software adapts to Level's language, relationships, criteria, approvals, capital

posture, and operating rhythm. Every signal the team processes, every memo the agents draft, every decision Levin approves or passes on, gets recorded in an owned context layer. The system gets sharper every week — and that compounding context is the asset.

### Three layers create the value.

- **An owned context layer** — Level's accounts, contacts, sites, permits, RFQs, bids, prior decisions, outcomes, and relationship memory, structured so software can actually use them.
- **An agentic workflow layer** — six specialist agents (Scout, Permit, Memory, Qualifier, Bid, Brief) that detect signals, pull context, score against Level's criteria, draft memos, and queue decisions for the right human.
- **An executive operating rhythm** — a weekly brief built from live system state, not chased status updates. Exception-first. Decision-ready.

**Level should not rent generic intelligence when it can build the operating memory that makes its own team sharper.**

The recommended first step is a focused **90-day pilot**: consolidate active opportunities, accounts, contacts, and RFQs into the context layer; ship Level-specific qualification scoring; stand up the RFQ and go/no-go command center; route approvals and next actions; and produce the first weekly executive intelligence brief from live data. The near-term outcome is cleaner visibility and faster decisions. The long-term outcome is a proprietary operating memory that improves every time Level reviews a signal, pursues an opportunity, wins a job, passes on a poor-fit project, or learns from execution.

---

## 2 • The context advantage

### WHY NOTHING ON THE MARKET DOES THIS FOR LEVEL

Every meaningful construction firm sits on a quiet asset: the cumulative judgment of its operators. For Level, that's the memory of who Carter is, why the Plano HGI bid won, what made the Frisco TI margin survive a tight schedule, which jurisdictions bleed weeks in plan review, and which brand operators are easy to bid against. That asset lives — today — in inboxes, spreadsheets, threads, RFQ folders, and people's heads.

Software can't use any of it. Search doesn't find it. New hires don't inherit it. Generic CRMs and project tools store records but not the relationships *between* records. That's the gap that limits

everything else: when a permit lands or a broker forwards a tip, Level's team can't ask *"what's our history with this owner, what similar pursuit did we win, which jurisdiction will slow this down, and who on our side should own the next move?"* — at least not in seconds, with structured context attached.

## What "context" means here

The Level Intelligence OS context layer is not a CRM, not a data lake, and not a document store. It's an **operating-grade graph of Level's business**: accounts, contacts, sites, parcels, markets, jurisdictions, permits, RFQs, bids, projects, vendors, capital sources, signals, decisions, approvals, and outcomes — all linked, all timestamped, all auditable.

That graph turns ten years of operating judgment into something an agent can query and a leader can see. Three properties make it valuable:

### OWNED

**100%**

Level's data, schema, and workflows — exportable, no vendor lock-in.

### COMPOUNDING

**↑ wk**

Every signal, score, override, and outcome makes the next decision sharper.

### QUERYABLE

**∞**

Natural-language access — Levin asks, the system answers from live records.

## The questions Level can finally answer

The context layer makes operating questions answerable on demand:

- "Show me Houston restaurant opportunities with strong relationship fit."
- "Which RFQs are due in the next 10 days and which are still missing documents?"
- "Summarize our history with this owner before the call at 3."
- "Which jurisdictions delayed us most often, and by how many days?"
- "Which opportunities should we pass on and why?"
- "What did we learn from lost bids this quarter?"
- "What changed since last week's executive brief?"

Today every one of those questions costs hours, three Slack threads, and a meeting. With the context layer in place, each is one query — answered with sourcing, scoring, and the relevant prior pursuit attached.

THE ASSET

Level's operating judgment is the asset. The context layer is the form that makes it usable, defensible, and durable — not stored in any one person's head, not held hostage by any one vendor, not erased when a key relationship leaves.

### 3 • Where Mercator fits — and where Level Intelligence OS goes further

Mercator is a useful category reference. It proves construction firms benefit from early project intelligence — surfaced permits, planning movement, ownership shifts, development signals. That category is real, and a Level Intelligence OS deployment brings that capability natively. But Mercator stops at the signal. Level needs the operating layer that takes a signal and turns it into a qualified, owned, routed, and converted pursuit.

CATEGORY REFERENCE

**Mercator-class intelligence**

- Surfaces public project signals (permits, planning, ownership).
- Generic feed with broad market coverage.
- Read-only — a data product Level subscribes to.
- No Level-specific qualification, no relationship memory.
- No RFQ workflow, no go/no-go, no approval queue.
- No executive operating rhythm tied to Level's calendar.

LEVEL INTELLIGENCE OS

**Owned operating layer**

- Brings the early-signal capability natively, scoped to Level's markets.
- Plus Level-specific qualification against your own criteria and history.
- Plus relationship memory — every account, contact, prior pursuit.
- Plus the RFQ and go/no-go command center on top of those signals.
- Plus the approval queue, audit trail, and executive brief.
- Plus the compounding context layer — owned and proprietary.

**We bring what Mercator does. We then keep going — qualifying, routing, briefing, approving, and remembering — inside one system that Level owns.**

The point is not a feature-by-feature replacement. The point is positioning: Mercator-class intelligence is the *input* layer. The operating value lives in what happens next — and the firm that owns the operating layer keeps the upside. When a permit surfaces, the right outcome isn't a notification, it's a scored memo with relationship history, comparable prior pursuit, and an owner pre-assigned — ready for Levin to approve before the morning meeting.

## **Other use cases beyond signals**

Because the context layer is general-purpose, the same operating system handles workflows Mercator-class tools never touch:

- **RFQ triage and go/no-go memos** — 128-page packages summarized, requirements extracted, missing documents flagged, comparable bids surfaced, memo drafted, approver queued.
- **Permit and entitlement monitoring** — plan review queues watched across jurisdictions, comments parsed, response drafted, days-in-review compared to historical cycle time.
- **Relationship memory and warm-intro mapping** — when a broker forwards a Schaumburg lead, the system surfaces D. Tan (Carter, ex-AC Hotels, 2-hop, last touch 38d) as the warm intro path.
- **Weekly executive brief** — Monday morning, the system polls every agent, assembles changes, stucks, and decisions, and pushes a decision-ready brief instead of a status meeting.
- **Capital and finance context** — opportunity intelligence linked to financing milestones, draw events, and investor reporting where relevant.
- **Software rationalization** — evidence on which of Level's current tools are earning their seat versus creating friction.

---

## **4 • The agentic stack**

SIX SPECIALISTS THAT COMPLETELY CHANGE THE WORKFLOW GAME

Each agent is a specialist operator with a narrow job, bounded tools, defined inputs, and explicit authority. None of them are vague general assistants. Together they collapse the loop from *signal arrives* to *decision queued* from **weeks** to a **single morning** — without removing human judgment from anything that matters.

## Scout

SIGNAL • DETECT

Watches permits, planning agendas, ownership filings, broker tips, RFQs. Extracts entities. Dedupes against the pipeline. Classifies asset type, market, jurisdiction. Routes to the right downstream agent.

### REPLACES TODAY

BD inbox triage, Sara forwarding "can you look at this?", three Slack threads asking if anyone's seen the filing.

## Permit

JURISDICTION • MONITOR

Watches plan review queues across every jurisdiction Level operates in. Parses plan review comments. Tracks days in review against historical cycle time. Flags blockers that affect pursuit or schedule assumptions.

### REPLACES TODAY

The PM checking portals weekly, the moment a Friday email surfaces a comment that's been open for nine days.

## Memory

RELATIONSHIPS • CONTEXT

Indexes accounts, contacts, prior pursuits. Pulls relationship strength, last touch, win/loss history, warm-intro paths. Generates the briefing one minute before Levin walks into a call.

### REPLACES TODAY

Texting the team asking "what do we know about Carter?", scrolling through three years of email looking for the last touch.

## Qualifier

SCORE • LEVEL CRITERIA

Scores every opportunity against Level's ten qualification dimensions — geography, asset, relationship, margin, capacity, strategic value. Recommends pursue / watch / pass. Identifies the missing information that would lift confidence.

### REPLACES TODAY

The conflicting numbers Sara and Mark gave last week, the inconsistent pursuit criteria across the team, the lost-bid post-mortems with no comparable lookup.

## Bid

RFQ • GO/NO-GO

Ingests RFQ packages (the 128-page hospitality ones included). Extracts scope, deadlines, submission format, insurance, bonding, missing documents. Drafts the go/no-go memo. Compares to prior bids. Flags schedule and FF&E risk.

### REPLACES TODAY

Two days of someone reading the package, the late-night realization the ESA was never in the bundle, the bid submitted with one risk we forgot to price.

## Brief

EXECUTIVE • ASSEMBLE

Polls every agent every Monday at 7:00 AM. Assembles changes, stucks, and decisions into one exception-first executive brief. Pushes it to founders, BD, precon, ops, and capital. The status meeting is now a 6-minute read.

### REPLACES TODAY

The Monday status meeting, the chase emails the Friday before, the half-fresh board update that's stale by the time it lands.

## How they actually work together

A new permit lands for Cedar Park HGI at 09:14. Inside the next six minutes:

- 09:14 · Permit parses the filing — Travis County, 120-key, applicant Carter.
- 09:15 · Scout dedupes against the pipeline and classifies hospitality / new-build.
- 09:15 · Memory pulls Carter — strength 0.78, 3 prior bids, 2 wins, last touch 14 days.
- 09:16 · Qualifier scores 0.84 with Plano HGI as comparable (won at 0.81).
- 09:18 · Bid drafts the pursuit memo with two open questions for Carter.
- 09:20 · Approval queued for M. Levin · move to pursue.

On the traditional path, the same signal sits in an inbox for three days, gets forwarded on day five, sits with stale CRM data on day nine, accumulates conflicting numbers by day sixteen, hits an OOO loop by day twenty-three, and dies on day thirty-one when the RFQ deadline passes. That contrast is the *workflow game change* — and the walkthrough at /two-paths animates it scroll-by-scroll. The agent ops console at /agent-ops lets you trigger the same scenarios live.

---

## 5 · ROI — eight compounding value levers

The strongest ROI argument isn't labor savings, though those are real. It's that Level's **owned operating context** compounds — every signal, every score, every override, every outcome makes the next decision sharper. Most software depreciates after install. This appreciates.

Eight levers create measurable value. The first four show up inside 90 days. The remaining four compound over the first 12-24 months.

## Near-term levers (Q1 of operation)

| LEVER                                | MECHANISM                                                                                                                                                               | HOW IT'S MEASURED                                                               |
|--------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------|
| <b>Opportunity capture</b>           | Signals detected from permits, planning, ownership, broker tips, and franchise news — converted into structured opportunity records with score, owner, and next action. | Signals captured / week · % converted to opportunity · win rate by source       |
| <b>Pursuit discipline</b>            | Every opportunity scored against Level's criteria and compared to historical pursuits. Faster, more consistent go/no-go decisions. Fewer low-margin distractions.       | Time to go/no-go · % with complete score · win rate by score band               |
| <b>Administrative load reduction</b> | Agents handle first-pass cognitive work — summarizing, extracting, classifying, comparing, drafting, routing, reminding. Status meetings shrink to exceptions.          | Hours saved per RFQ · hours saved per weekly brief · reduction in overdue tasks |
| <b>RFQ and bid quality</b>           | RFQ packages parsed in minutes, missing documents flagged, comparable prior bids surfaced, go/no-go memo drafted. Level bids fewer, but bids better.                    | RFQs processed · avg review time · % with memo · win rate by RFQ type           |

## Compounding levers (12-24 months)

| LEVER                                      | MECHANISM                                                                                                                                                                  | HOW IT'S MEASURED                                                                               |
|--------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------|
| <b>Permit / entitlement risk reduction</b> | Permits monitored across jurisdictions. Comments parsed, responses drafted, blockers flagged. Earlier detection of delays, faster response, jurisdiction memory grows.     | Avg days in review<br>· % past expected cycle<br>· time from comment to response                |
| <b>Relationship memory</b>                 | Account timelines, contact histories, last touch, next touch, warm intros. Relationship knowledge no longer trapped in one person's head — durable across team changes.    | Stale relationship count<br>· follow-up completion<br>· opportunities sourced by relationship   |
| <b>Executive focus</b>                     | Weekly brief assembled from live data. Pending decisions surfaced. Stalled opportunities and risks flagged. Leadership time shifts from status gathering to judgment.      | Pending decisions by age<br>· stalled opportunity count<br>· brief usage / adoption             |
| <b>Enterprise data asset</b>               | Every signal, score, decision, relationship, RFQ, bid, permit, vendor outcome, win/loss, and margin result becomes structured memory. Proprietary, exportable, defensible. | Structured record count<br>· recommendation accuracy<br>· outcomes tied back to prior decisions |

### THREE SCENARIO SHAPES (TO BE POPULATED IN DISCOVERY)

**Scenario A — Operational efficiency only.** Value comes from administrative time savings, fewer missed deadlines, fewer status meetings, faster RFQ triage, faster permit follow-up. Justifies the pilot on its own.

**Scenario B — One better opportunity.** Value comes from identifying, qualifying, or winning one additional meaningful project per year that would otherwise have been missed or pursued too late. Typically dwarfs Scenario A by an order of magnitude.

**Scenario C — Enterprise operating asset.** Value comes from the system improving the quality and consistency of Level's decisions over time — better scoring, reusable memory, onboarding leverage, owned software optionality. This is the durable enterprise-value lever.

### What we recommend Level *not* claim

Don't sell the pilot on labor savings alone. Labor savings are real and easy to measure, but they undersell the asset. The larger value — and the more honest one — is **better opportunities, better decisions, less leakage, institutional memory, faster adaptation, and enterprise maturity**. That's what the context layer is for.

---

## 6 • Procore bridge strategy

WRAP FIRST. EARN THE RIGHT TO REPLACE.

Procore is a strong source of truth for active project execution. The first job is to structure Level Intelligence OS *around* that reality — adding the intelligence and executive layer Procore is not designed to be — rather than forcing a disruptive migration.

| PHASE                                     | POSITION                                                                                                                     | OUTCOME                                                            |
|-------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------|
| <b>1 • Read and reference</b>             | Treat Procore as the current execution source of truth. Pull or reference only the data needed for visibility and decisions. | Zero disruption to field or project teams.                         |
| <b>2 • Workflow overlay</b>               | Use Level Intelligence OS for signals, RFQs, go/no-go memos, relationship memory, approvals, and executive briefs.           | Operating rhythm improves without a platform migration.            |
| <b>3 • Evidence-based rationalization</b> | Measure which workflows are better handled inside Level's owned system and which should stay in Procore.                     | Software decisions based on observed value, not vendor narratives. |
| <b>4 • Selective replacement</b>          | Move modules away from monolithic software only when the owned layer is clearly better, safer, and cheaper.                  | Dependency reduces over time without risky rip-and-replace.        |

The operating principle: *earn the right to replace.*

## 7 • Human control and trust

The system prepares work. It does not silently make consequential decisions. Authority is explicit, audit trails are required, and Level controls the final move on anything that touches the outside world or commits resources.

| AUTHORITY LEVEL          | SYSTEM CAN DO                                                                           | HUMAN CONTROL                                |
|--------------------------|-----------------------------------------------------------------------------------------|----------------------------------------------|
| <b>Autonomous</b>        | Classify signals, summarize RFQs, dedupe contacts, draft briefs, extract requirements.  | Logged for review.                           |
| <b>Suggest</b>           | Recommend owner, next action, pursue/watch/pass, follow-up draft.                       | Human accepts, edits, or overrides.          |
| <b>Approval required</b> | External outreach, bid submission, vendor award, budget change, investor communication. | Explicit approval required before execution. |
| <b>Manual</b>            | Legal, contract, financing, final strategic decisions.                                  | Human executes; agents assist.               |

Every meaningful recommendation carries: what the system recommends, why, the evidence it used, what's missing, its confidence, the approver, the decision, and eventually the outcome. That audit trail is what lets the system learn from leadership overrides instead of repeating them.

## 8 • The 90-day pilot

DESIGN, BUILD, TEST, OPERATIONALIZE

### Objective

Prove that Level Intelligence OS can improve opportunity visibility, qualification discipline, RFQ coordination, executive reporting, and institutional memory within 90 days — and produce a phase-two roadmap that's grounded in evidence rather than assumption.

### Scope (v1)

- Active opportunity pipeline.
- Account and contact relationship layer.
- Signal inbox (Scout + Permit feeding the context layer).
- Level-specific qualification scoring.
- RFQ intake and go/no-go memo workflow.
- Approval and next-action queue.
- Weekly executive brief, generated from live system data.
- Outcome capture and feedback loop.

### Timeline

|                                                                                                                                                     |                                                                                                                                                             |                                                                                                                                                |                                                                                                                                            |
|-----------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|
| <div>WEEKS 1-2</div> <div>Foundation</div> <div>Data model, roles, active opportunities, account/contact views, executive dashboard skeleton.</div> | <div>WEEKS 3-5</div> <div>Intelligence</div> <div>Signal inbox, opportunity summaries, qualification score, relationship briefs, first weekly digest.</div> | <div>WEEKS 6-8</div> <div>Workflow</div> <div>RFQ command center, deadline extraction, go/no-go memo, approval queue, bid task tracking.</div> | <div>WEEKS 9-12</div> <div>Operating rhythm</div> <div>Weekly executive brief, outcome tracking, adoption review, phase-two roadmap.</div> |
|-----------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|

## Success metrics

Pilot success is measured by operating usefulness, not novelty:

- 100% of active pursuits visible in the system, each with owner, stage, source, score, next action, and age.
- At least 10 useful signals captured, reviewed, and routed during the pilot.
- RFQs tracked with deadlines, requirements, missing-document flags, and go/no-go status.
- Weekly executive brief generated from live system data, used in the leadership cadence.
- Leadership identifies faster decisions, fewer blind spots, or clearer accountability than before the pilot.
- Phase-two roadmap approved, adjusted, or rejected based on evidence collected during the pilot.

## Phase-two options

Expansion is sequenced based on where Level sees the highest value during the pilot:

- Permit and entitlement monitoring at jurisdiction depth.
- Email, calendar, and CRM integration into the context layer.
- Procore read-only bridge.
- Vendor and subcontractor performance memory.
- Project handoff from pursuit to execution.
- Finance and capital workflow.
- Investor reporting support.
- Hospitality pipeline tracking.
- Software rationalization and existing-tool usage audit.

---

## 9 • Commercial principles and ownership

The final commercial structure will be defined after discovery. The operating principles, however, should be clear up front.

### Engagement

90-day pilot to design, build, test, and operationalize Level Intelligence OS v1.

## Pilot deliverables

- Level data model and workflow map.
- Active opportunity / account / contact database.
- Signal inbox and triage workflow.
- Qualification scoring framework.
- RFQ intake and go/no-go memo workflow.
- Approval queue.
- Weekly executive brief.
- Pilot success report and phase-two roadmap.

## Ownership

- **Level owns its data**, operating memory, domain model, and workflow definitions.
- Level retains export rights for pilot data and generated records.
- Code ownership and licensing defined in the final agreement.
- AI model providers remain replaceable infrastructure — not a lock-in surface.
- The strategic asset is Level's operating judgment, made structured and reusable.

## Data and confidentiality

- Level data is confidential.
- Role-based access and audit logs required.
- No external reuse of Level data without written permission.
- Sensitive actions require human approval.
- Data export available at all times.

---

# 10 • Discovery questions

A short, workflow-grounded discovery sharpens the pilot. The questions below are the ones we'd work through in the first two weeks — not abstract AI strategy questions, but the practical anchors that determine where Level Intelligence OS sits inside Level's actual operating rhythm.

- Where do new opportunities originate today — and which sources matter most?

- Who owns opportunity qualification, and where do go/no-go decisions get made?
  - What makes a Level opportunity attractive or unattractive — by asset, market, owner, capacity?
  - Where do RFQs live today, and how is the package read, scoped, and decided?
  - What fields does leadership need to see every week — and which ones are stale today?
  - Which relationships, markets, and asset types matter most over the next 12 months?
  - Which systems are in use today — Procore, CRM, email, calendar, accounting, storage?
  - Which Procore modules are true sources of truth, and which are mainly used for visibility?
  - What data can be imported safely for the pilot?
  - Which actions require explicit approval, and from whom?
  - What would make the pilot obviously worth continuing into phase two?
- 

## 11 • Closing position

Level Intelligence OS should be read as a serious operating asset, not an AI experiment. The first version gives Level one place to see opportunities, understand fit, track RFQs, preserve relationship memory, and route decisions. The long-term version becomes Level's private construction intelligence layer — a system that learns from Level's signals, relationships, decisions, wins, losses, and operating judgment, and gets sharper every quarter.

**Mercator finds opportunities. Level Intelligence OS finds, qualifies, routes, and converts opportunities using Level's own operating judgment — on top of an owned context layer that compounds. The first is a feed Level subscribes to. The second is an asset Level owns.**

The recommended next step is a 90-day pilot, scoped tightly enough to ship and operationalize inside one quarter, generous enough to prove the compounding loop with real Level data. From there, phase two is decided from evidence, not assumption.

---

## 12 • Companion artifacts

This document is one of three artifacts in the Level Intelligence OS package. The walkthrough animations make the operating model visible. Use them together.

ANIMATION • SCROLL SITE

## Two Paths

Same morning, two paths. A scroll-driven walkthrough of the 31-days-vs-six-minutes contrast across the full pipeline.

LIVE CONSOLE

## Agent Ops

A live demo of the agent network. Trigger four scenarios — permit, RFQ, broker tip, weekly brief — and watch the agents work in parallel.

A PDF of this proposition is available via the [download link in the header](#) — share it with stakeholders who'd rather read offline.